

Bingqi Shang

 [Website](#) |  [Google Scholar](#) |  [GitHub](#) |  bshang@msu.edu

RESEARCH INTERESTS

Trustworthy Machine Learning: Machine Unlearning, Alignment & RLHF, Adversarial Machine Learning

Generative AI: Large-language Model, Multi-modal Language Model, Reasoning, Post-training

EDUCATION

Michigan State University (MSU)

Ph.D., Computer Science

Sep. 2025 - Present

Advisor: [Prof. Sijia Liu](#)

Northwestern University (NU)

M.S., Computer Science

Sep. 2023 - Jun. 2025

Advisors: [Prof. Qi Zhu](#) and [Prof. Xiao Wang](#)

Tongji University

B.E., Software Engineering

Sep. 2019 - Jun. 2023

[School of Computer Science and Technology](#)

RESEARCH EXPERIENCE

Backdooring Machine Unlearning via Attention Sinks

Apr. 2025 - Oct. 2025

Supervisor: [Prof. Sijia Liu](#) (MSU)

- Exposed a vulnerability in LLM unlearning: attention-sink-guided backdoors can **reactivate** forgotten behaviors.
- Exploring applications in unlearned models where backdoor triggers can selectively recover forgotten knowledge.
- **Publication:** [2].

Privacy-Preserving Tuning for Large Models

Dec. 2023 - Mar. 2025

Supervisors: [Prof. Qi Zhu](#) (NU), [Prof. Xiao Wang](#) (NU)

- Developed Split Adaptation (SA) to ensure **data privacy** during adaptation of pre-trained Vision Transformers (ViTs), utilizing bi-level noise injection for privacy-preserving downstream tasks without data sharing.
- Protected **model privacy** by sharing only a low-bit quantized frontend of the ViT, preventing model leakage and ensuring secure adaptation.
- **Publication:** [3].

PROFESSIONAL EXPERIENCE

SAP Shanghai, China

Jun. 2022 - Mar. 2023

Cloud Developer Intern, Supervisor: April Qi

Project: Cloud Provider Exporter in Go on Kubernetes for AWS, Azure, and GCP, using Prometheus and Grafana.

PUBLICATIONS

* indicates equal contribution

[1] Yiwei Chen*, **Bingqi Shang***, Sijia Liu. **Who Built This Model? Tracing LLM Lineage via Spectral Fingerprints in Weight Space.** *COLM'2026*.

[2] **Bingqi Shang***, Yiwei Chen*, Yihua Zhang, Bingquan Shen, Sijia Liu. **Forgetting to Forget: Attention Sink as A Gateway for Backdooring LLM Unlearning.** *COLM'2026*.

[3] Lixu Wang*, **Bingqi Shang***, Yi Li, Payal Mohapatra, Wei Dong, Xiao Wang, Qi Zhu. **Split Adaptation for Pre-trained Vision Transformers.** *CVPR'2025*.

[4] Yancheng Huang*, Changsheng Wang*, Chongyu Fan, Yicheng Lang, **Bingqi Shang**, Yang Zhang, Mingyi Hong, Qing Qu, Alvaro Velasquez, Sijia Liu. **Subspace Control: Turning Constrained Model Steering into Controllable Spectral Optimization.** *COLM'2026*.

[5] Zijie Liu, Jinhao Duan, **Bingqi Shang**, Xinming An, Sijia Liu, Tianlong Chen. **When Retain Constraints Conflict: Mitigating Forget-Retain Interference in Tabular Data.** *EMNLP'2026 Findings*.

HONORS

- Shanghai Outstanding Graduate Award

2023

- Outstanding Undergraduate Dissertation Award of Tongji University 2023
- **National Scholarship** (Top 0.2%, highest undergraduate honor in China) 2020

SERVICES

Workshop Co-organizer: AdvML-Frontiers × CoTMA [COLM'26]

Journal Reviewer: IEEE TSP, IEEE TAES

Conference Reviewer: NeurIPS 2026, CVPR 2026, ICASSP 2026–2027

PERSONAL INTERESTS

Photography 2019 - Present

PROFESSIONAL SKILLS

Programming Languages: Python, Go, C++, Java, Rust, L^AT_EX, HTML, CSS, JavaScript

Machine Learning Systems: PyTorch, Transformers, Huggingface, W&B, OpenCV, Scikit-learn